

PRACTICAL NO.-4

Aim: Run a java program based on parallel programming to implement the concept of Map Reduce Paradigm.

DESCRIPTION:-

The **WordCount** program is the classic example for learning **Hadoop MapReduce**. It counts the frequency of each word in a given input file. This practical demonstrates how to:

1. Prepare input data in **HDFS**.
2. Write **Mapper, Reducer, and Driver** classes in Java.
3. Compile and create a **JAR file**.
4. Run the MapReduce job on Hadoop through **PuTTY**.
5. View results stored in HDFS.

This experiment helps understand the **basic flow of MapReduce** in Hadoop.

STEPS OF EXECUTION: -

Step 1: Prepare Input Data

```
echo "Hadoop is big data Hadoop is Java" > sample.txt
hdfs dfs -mkdir /input
hdfs dfs -put sample.txt /input/
hdfs dfs -ls /input
```

Step 2: Create Java Files

WordMapper.java

```
import java.io.IOException;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.LongWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Mapper;
public class WordMapper extends Mapper<LongWritable, Text, Text, IntWritable> {
    private final static IntWritable one = new IntWritable(1);
    private Text word = new Text();
    public void map(LongWritable key, Text value, Context context)
        throws IOException, InterruptedException {
        String line = value.toString();
        for (String token : line.split("\s+")) {
            word.set(token);
            context.write(word, one);
        }
    }
}
```

```
}
```

WordReducer.java

```
import java.io.IOException;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Reducer;

public class WordReducer extends Reducer<Text, IntWritable, Text, IntWritable> {
    public void reduce(Text key, Iterable<IntWritable> values, Context context)
        throws IOException, InterruptedException {
        int sum = 0;
        for (IntWritable val : values) {
            sum += val.get();
        }
        context.write(key, new IntWritable(sum));
    }
}
```

WordCountDriver.java

```
import org.apache.hadoop.conf.Configuration;
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Job;
import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;

public class WordCountDriver {
    public static void main(String[] args) throws Exception {
        Configuration conf = new Configuration();
        Job job = Job.getInstance(conf, "word count");

        job.setJarByClass(WordCountDriver.class);
        job.setMapperClass(WordMapper.class);
        job.setReducerClass(WordReducer.class);
        job.setOutputKeyClass(Text.class);
        job.setOutputValueClass(IntWritable.class);
        FileInputFormat.addInputPath(job, new Path(args[0]));
        FileOutputFormat.setOutputPath(job, new Path(args[1]));

        System.exit(job.waitForCompletion(true) ? 0 : 1);
    }
}
```

Step 3: Compile Java Files

```
mkdir wordcount_classes
```

```
hadoop com.sun.tools.javac.Main -d wordcount_classes WordMapper.java WordReducer.java  
WordCountDriver.java
```

☞ If error occurs, use:

```
javac -cp `hadoop classpath` -d wordcount_classes WordMapper.java WordReducer.java  
WordCountDriver.java
```

Step 4: Create JAR File

```
jar -cvf wordcount.jar -C wordcount_classes/ .
```

Step 5: Run MapReduce Job

```
hdfs dfs -rm -r /output  
hadoop jar wordcount.jar WordCountDriver /input /output
```

Step 6: View Output

```
hdfs dfs -ls /output  
hdfs dfs -cat /output/part-r-00000
```

 Expected Output Example:

```
Hadoop 2  
Java 1  
big 1  
data 1  
is 2
```

```
[maria_dev@sandbox-hdp ~]$ hdfs dfs -ls /input  
^C[maria_dev@sandbox-hdp ~]$ ls  
sample.txt u.data wordcount_classes  
[maria_dev@sandbox-hdp ~]$ vi WordMapper.java  
[maria_dev@sandbox-hdp ~]$ ls  
sample.txt u.data wordcount_classes WordMapper.java  
[maria_dev@sandbox-hdp ~]$ vi WordReducer.java  
[maria_dev@sandbox-hdp ~]$ vi WordCountDriver.java  
[maria_dev@sandbox-hdp ~]$ mkdir wordcount_classes  
mkdir: cannot create directory 'wordcount_classes': File exists  
[maria_dev@sandbox-hdp ~]$ hadoop com.sun.tools.javac.Main -d wordcount_classes  
WordMapper.java WordReducer.java  
Error: Could not find or load main class com.sun.tools.javac.Main  
[maria_dev@sandbox-hdp ~]$ WordCountDriver.java  
-bash: WordCountDriver.java: command not found  
[maria_dev@sandbox-hdp ~]$ javac -cp `hadoop classpath` -d wordcount_classes Wo  
rdMapper.java WordReducer.java  
  
WordMapper.java:1: error: class, interface, or enum expected  
import java.io.IOException;  
^  
WordMapper.java:12: error: illegal escape character  
    for (String token : line.split("\s+")) {  
        ^  
WordReducer.java:1: error: class, interface, or enum expected  
import java.io.IOException;  
^  
3 errors  
[maria_dev@sandbox-hdp ~]$ WordCountDriver.java  
-bash: WordCountDriver.java: command not found  
[maria_dev@sandbox-hdp ~]$ javac -cp `hadoop classpath` -d wordcount_classes Wo  
rdMapper.java WordReducer.java  
^[[D[[WordMapper.java:1: error: class, interface, or enum expected  
import java.io.IOException;
```

```

3 errors
[maria_dev@sandbox-hdp ~]$ WordCountDriver.java
-bash: WordCountDriver.java: command not found
[maria_dev@sandbox-hdp ~]$ javac -cp `hadoop classpath` -d wordcount_classes WordMapper.java WordReducer.java
"ID" [WordMapper.java]: error: class, interface, or enum expected
import java.io.IOException;
^
WordMapper.java:12: error: illegal escape character
    for (String token : line.split("\n+")) {
                           ^
WordReducer.java:1: error: class, interface, or enum expected
import java.io.IOException;
^
3 errors
WordMapper.java WordReducer.java$acc -cp `hadoop classpath` -d Wordcount_classes WordCountDriver.java

WordMapper.java:1: error: class, interface, or enum expected
import java.io.IOException;
^
WordMapper.java:12: error: illegal escape character
    for (String token : line.split("\n+")) {
                           ^
WordReducer.java:1: error: class, interface, or enum expected
import java.io.IOException;
^
3 errors
[maria_dev@sandbox-hdp ~]$ WordCountDriver.java
-bash: WordCountDriver.java: command not found
[maria_dev@sandbox-hdp ~]$ vi WordReducer.java
[maria_dev@sandbox-hdp ~]$ vi WordCountDriver.java
[maria_dev@sandbox-hdp ~]$ vi WordMapper.java
[maria_dev@sandbox-hdp ~]$ javac -cp `hadoop classpath` -d Wordcount_classes WordMapper.java WordReducer.java
WordMapper.java:12: error: illegal escape character
    for (String token : line.split("\n+")) {
                           ^
1 error
[maria_dev@sandbox-hdp ~]$ vi WordMapper.java
[maria_dev@sandbox-hdp ~]$ ^C
[maria_dev@sandbox-hdp ~]$ vi WordMapper.java
[maria_dev@sandbox-hdp ~]$ vi WordMapper.java
[maria_dev@sandbox-hdp ~]$ javac -cp `hadoop classpath` -d wordcount_classes WordMapper.java WordReducer.java
WordReducer.java:9: error: incompatible types: Object cannot be converted to IntWritable
    for (IntWritable val : values) {
                           ^
Note: Some input files use unchecked or unsafe operations.
Note: Recompile with -Xlint:unchecked for details.
1 error
[maria_dev@sandbox-hdp ~]$ javac -cp `hadoop classpath` -d wordcount_classes WordMapper.java WordReducer.java WordCountDriver.java
WordReducer.java:9: error: incompatible types: Object cannot be converted to IntWritable
    for (IntWritable val : values) {
                           ^
Note: Some input files use unchecked or unsafe operations.

Note: Some input files use unchecked or unsafe operations.
Note: Recompile with -Xlint:unchecked for details.
1 error
[maria_dev@sandbox-hdp ~]$ ^C
[maria_dev@sandbox-hdp ~]$ vi WordReducer.java
[maria_dev@sandbox-hdp ~]$ javac -cp `hadoop classpath` -d wordcount_classes WordMapper.java WordReducer.java WordCountDriver.java
Note: WordMapper.java uses unchecked or unsafe operations.
Note: Recompile with -Xlint:unchecked for details.
[maria_dev@sandbox-hdp ~]$ ^C
[maria_dev@sandbox-hdp ~]$ jar -cvf wordcount.jar -C wordcount_classes/
added manifest
adding: WordMapper.class(in = 1275) (out= 690)(deflated 46%)
adding: WordReducer.class(in = 1592) (out= 663)(deflated 58%)
adding: WordCountDriver.class(in = 1364) (out= 752)(deflated 44%)
[maria_dev@sandbox-hdp ~]$ 
[maria_dev@sandbox-hdp ~]$ hdfs dfs -rm -r /output

rm: '/output': No such file or directory
[maria_dev@sandbox-hdp ~]$ 
[maria_dev@sandbox-hdp ~]$ ^C
[maria_dev@sandbox-hdp ~]$ hadoop jar wordcount.jar WordCountDriver /input /output

25/09/01 04:57:03 INFO client.RMProxy: Connecting to ResourceManager at sandbox-hdp.hortonworks.com/172.18.0.2:8032
25/09/01 04:57:04 INFO client.AMRMProxy: Connecting to Application History server at sandbox-hdp.hortonworks.com/172.18.0.2:18200

```

```
File System Counters
  FILE: Number of bytes read=82
  FILE: Number of bytes written=305979
  FILE: Number of read operations=0
  FILE: Number of large read operations=0
  FILE: Number of write operations=0
  HDFS: Number of bytes read=155
  HDFS: Number of bytes written=34
  HDFS: Number of read operations=6
  HDFS: Number of large read operations=0
  HDFS: Number of write operations=2
```

CONCLUSION: -

In this practical, we implemented and executed the **Hadoop WordCount program** on a singlenode Hadoop cluster using PuTTY. We successfully:

- Uploaded input data into HDFS,
- Compiled Java Mapper, Reducer, and Driver classes,
- Packaged them into a JAR,
- Executed the MapReduce job, and
- Verified the word frequency results from HDFS.

This experiment demonstrates the **basic working of MapReduce** in Hadoop.